

COM6002 Big Data Management

Data structure

Objectives

- Understand the meaning of Big-O
- Apply various data structures in Python
- Apply parallel processing in Python
- What is not included in this chapter?
 - Most of the theories and concepts behind the data structures
 - We focus on usage (coding!) here

LeetCode

- A free online platform
- Has a great database of questions for coding
 - Some questions are real job interview questions from big companies like Microsoft
 - Most with answers and some with tutorials
- Cover many basic concepts in programming and data structure
- Not just to achieve the functional goal, but ensure your program is **fast enough** to run within the time limit

Motivating example

- Check experiment 1
- Key observation
 - The same code segment but different data structure
 - Lead to a significant difference in running time!

How do we analyze the performance?


- In computer science, we usually use complexity
- Big O notation
 - Measure the order of the running cost (time / space etc.)
- Example 1:
 - For each student in the class
 - Mark his/her assignment
- The above task has a complexity of $O(n)$ where n is the number of students in the class
- The teacher's workload grows linearly with n

Big O notation examples

- Example 2:
 - For each student in the class
 - The student writes a peer evaluation report to every other student in the class
 - For each peer evaluation report
 - The teacher marks the report
- The above marking task has a complexity of $O(n^2)$ where n is the number of students in the class
- The teacher's workload grows linearly with n^2

Big O notation examples

- Example 3:
 - While there are more than one player left
 - Every two players play a game: the winner stays and the loser leaves
 - ** Many two-player games can happen at the same time; we call it a round
- Number of rounds is $O(\lg n)$
- Number of games is $O(n)$
- Where n is the number of players



Big O notation examples

- Example 4:
 - Find any ONE student in the class
- Cost is $O(1)$
- $O(1)$ means it is constant time regardless of how big n is
 - Where n is the number of students in the class

Big O rules


- If the cost is $3n^2 + 4n + 100 \rightarrow O(n^2)$
 - We only care about the highest order term
 - We don't care about the coefficient
- We usually consider the average case or the worst case
 - Example:
 - For each student in the class
 - If the student is Tom, bingo and quit; otherwise, keep looking
 - Best case: the first student is a hit, $O(1)$
 - Worst case: $O(n)$
 - Average: $O(n)$ --- We need to check half of the list on average

Python data structures

- Native
 - List
 - Set
 - Dict
- Commonly used packages
 - Deque
 - Heap
- Check the provided sample codes to see how we can use them in programming

Understanding the data structure

- We use a (physical) notebook to explain the concepts



- Python List
 - The first data element goes to page 1
 - The second data element goes to page 2
 - ... and so on

Cost analysis of Python List

- Adding a new element in the middle of the list?
 - $O(n)$
 - Why?
 - You need to move elements backwards
 - Insert to page 2
 - The old page 2 moves to page 3
 - The old page 3 moves to page 4
 - ... and so on
- For all discussions on the costs, n refers to number of elements in the data structure

Cost analysis of Python List

- Getting the i -th element [also called random access]
 - $O(1)$
 - Easy. Go to page i directly
- Delete the i -th element
 - $O(n)$
- Adding element at the back (push / append)
 - $O(1)$
- Deleting the last element (pop)
 - $O(1)$
- Query: is x in the data structure?
 - $O(n)$

Understanding set

- Like mathematical set
 - No duplicated element
 - Example
 - If we put $[1, 3, 3, 5, 5, 5]$ into a set, it keeps $\{1, 3, 5\}$
 - There is no order between the data
 - When you loop the elements in the set, there is no guaranteed order of data
- Behind the scene, set is implemented as a **hash table**

Mechanism of set

- Assume the notebook has 100 pages
- We use a hash function to convert an element to $[1, 100]$ or $[0, 99]$
 - Note: say 0 means page 100
 - In reality, the hash function also needs to consider string / float / boolean data etc. as input
- Example:
 - Hash function $h(x) = x \% 100$
 - % is the modulo operator
 - Do the division and take the remainder
 - Examples: $16\%5 = 1$, $13\%7 = 6$, $23\%4 = 3$
 - Element $33967 \% 100$ becomes 67
 - Then, we write the data 33967 on page 67


Collision

- What happens if we want to put two elements with the same hash value to the same page
 - Example elements:
 - 339⁶⁷ and 287⁶⁷
- There are different collision resolve mechanisms
- In our notebook case, just imagine we write both data elements on the same page
 - So, in the worst case, we write all the data on the same page ➔ searching becomes very difficult / slow: $O(n)$

Cost analysis

- Adding a new element
 - Average: $O(1)$; Worst: $O(n)$
- Deleting an element
 - Average: $O(1)$; Worst: $O(n)$
- Query: is x in the data structure?
 - Average: $O(1)$; Worst: $O(n)$
- Q: Is set better than list?
 - Think about how you add a new element
 - Don't just look at complexity

Understanding dict

- Similar to set
- Keep key-value pairs
 - In the form of {key: value}
- Example:
 - {1: 30, "a": 25, 30: "c"}
 - Similar to a set of {1, "a", 30}; Then we keep additional information about the data element

Cost analysis

- Adding a new key
 - Average: $O(1)$; Worst: $O(n)$
- Deleting a key
 - Average: $O(1)$; Worst: $O(n)$
- Query: is x in a key of dict?
 - Average: $O(1)$; Worst: $O(n)$

Understanding deque

- Data element is written on a random page
- We also write the the page numbers of the previous data element and the next data element
- If there is no sibling page (this is the last element; there is no next page), we write “None” (or null)
- We remember the first page number (and the last page for a doubly linked list)
- Behind the scene, deque is a **doubly linked list**

Deque examples

- At the beginning, there is no data element

First Page:
None
Last Page:
None

- Say, an element 13 is added
 - We find an empty page (say, 52) and write the data element on the page

First Page:
52
Last Page:
52

Data: 13
Prev: None
Next: None

Page 52

Deque examples

- Say, an element 26 is added
 - We find an empty page (say, 78) and write the data element on the page

First Page:
52
Last Page:
78

Data: 13
Prev: None
Next: 78

Page 52

Data: 26
Prev: 52
Next: None

Page 78

- Adding 30 on page 11

First Page:
52
Last Page:
11

Data: 30
Prev: 78
Next: None

Page 11

Data: 13
Prev: None
Next: 78

Page 52

Data: 26
Prev: 52
Next: 11


Page 78

Simplified visual for linked list

- Adding 13



- Adding 26



- Adding 30



- Pros:

- Easy handling of adding / deleting data in the middle
 - Does not require moving data forward / backward

- Cons:

- Not easy for random access

Cost analysis

- Getting the i -th element [also called random access]
 - $O(n)$
- Adding a new element (assume we know where to add)
 - $O(1)$
- Deleting an element (assume we know where to delete)
 - $O(1)$
- Query: is x in the data structure?
 - $O(n)$


Understanding Heap

- Also called priority queue
 - Very efficient in finding the smallest (or largest) element in a queue
 - Allow adding new data to the queue
- Consider this problem:
 - Find the smallest number --- $O(n)$
 - Find again the 2nd smallest number in the remaining data --- $O(n)$
 - Find again the 3rd smallest number --- $O(n)$
 - ...
 - Outputting the data from smallest to largest takes $O(n^2)$

Can we reuse some previous information to reduce the cost?


Heap structure

- View the data list as a tree-like structure
- Converting from a list to a heap structure is called “**heapify**”




How do we find the smallest value?

- Like a tournament
 - Smallest of the smallest




Finding the smallest in a group

- 9 is the smallest




11	31	54	9	18	77	81
----	----	----	---	----	----	----

Swapping the positions




Check the other group

- 54 is the smallest. Keep the positions of all




Continue

- 9 is the smallest. Swap the positions




Result

- The root node (the node at the top level) is guaranteed to be the smallest
 - Q: Can you prove it?




Heappop

- Heappop:
 - Remove the smallest element from the queue
- As a list, it is efficient to remove the last element
- We move the last the element to the first




Heappop maintenance

- The heap structure has a property that the smallest node should be placed as the parent
- 11 should be swapped with 81




Heappop maintenance

- The same applies to the left side child tree
 - 18 swaps with 81
- We do not need to examine the right child
 - The swap at the higher level only affects the left child



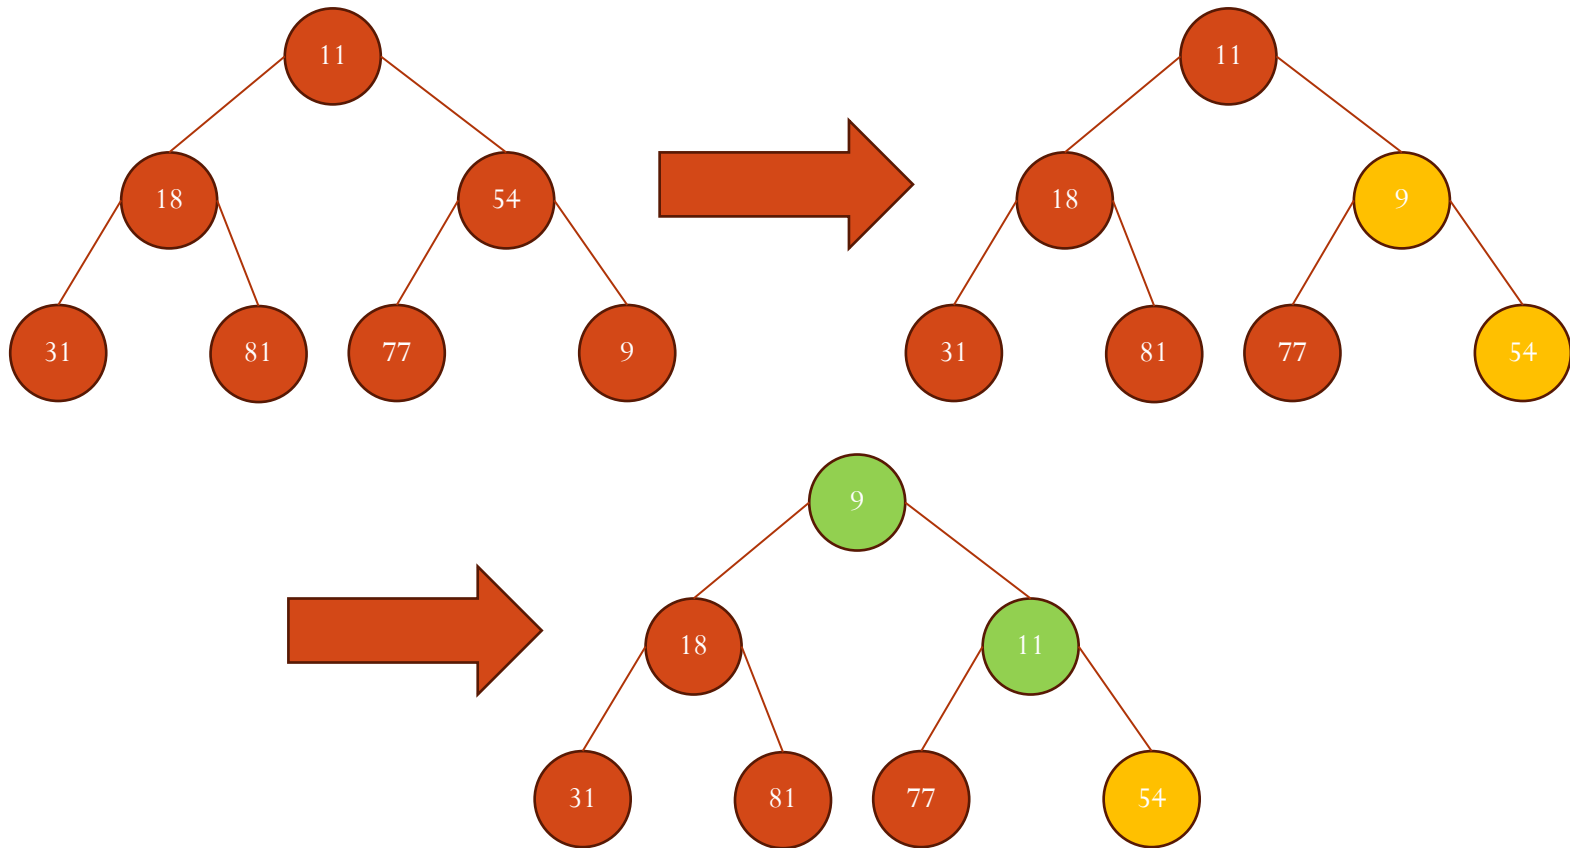
Heappop result



Heappush:


Adding new element to the queue

- Say adding 9
 - Update from the bottom to the top



Heappush: Another example

- Say adding 29
 - Update from the bottom to the top



Cost analysis

- Heapify
 - $O(n)$
- Heappush
 - $O(\lg(n))$
- Heappop
 - $O(\lg(n))$

More information [self-study]

- More about the complexity information of Python data structures
 - <https://wiki.python.org/moin/TimeComplexity>
- More information about Heap
 - https://en.wikipedia.org/wiki/Binary_heap

Example: Keys and Rooms

- <https://leetcode.com/problems/keys-and-rooms/>
- Your [non-coding] task:
 - Try to describe a solution to your classmate
 - Make sure it can be easily understood
 - Listen to your classmate's solution
 - Try to understand what it is doing
 - Reason about the correctness of the solution
 - Imagine someone is stupid and will just follow your instructions to work. Will he/she get the correct answer?
 - Computer is stupid and will just follow your instructions

Step #1 – the basic visiting logic

1. Visit room 0
2. Add every room it can open to a queue / list (or whatever structure that is appropriate) “q”
3. Pick any room from “q”
4. Go back to step 2 until “q” is empty

Step #2 – count how many rooms are visited

1. Visit room 0
2. Keep a list / queue (or whatever structure that is appropriate) “visited” about visited rooms
3. Add room 0 to “visited”
4. Add every room it can open to a queue / list (or whatever structure that is appropriate) “q”
5. Add all these rooms to “visited”
6. Pick any room from “q”
7. Go back to step 4 until “q” is empty

Step #3 – avoid infinite loop

1. Visit room 0
2. Keep a list / queue (or whatever structure that is appropriate) “visited” about visited rooms
3. Add room 0 to “visited”
4. Add every room it can open **and is not in “visited”** to a queue / list (or whatever structure that is appropriate) “q”
5. Add all these rooms to “visited”
6. Pick any room from “q”
7. Go back to step 4 until “q” is empty

Step #4 – determine the most efficient structure

1. Visit room 0
2. Keep a **set** “visited” about visited rooms
3. Add room 0 to “visited”
4. Add every room it can open **and is not in “visited”** to a list “q”
5. Add all these rooms to “visited”
6. Pick **the last room** from “q” (if any room can do, we simply pick the last room for **efficiency purpose**)
7. Go back to step 4 until “q” is empty

Example:


- <https://leetcode.com/problems/path-with-maximum-probability/>
 - This example is more advanced
- Q: What is your approach to solve the problem?

Sub-problem: edge access efficiency

- The edges are given as a list
 - How can we determine if there is a path from x to y ?
 - Complexity: $O(e)$
 - Where e is number of edges. Number of edges is of $O(n^2)$ where n is number of nodes
 - Sample method
 - Convert it to dict of dict
- ```
path: dict[int, dict[int, float]] = {}
```
- $\text{path}[u][v]$  = probability of the edge from  $u$  to  $v$
  - Q: Complexity to determine if there is a path from  $x$  to  $y$ ?

# Solution logic


- Idea:
  - Keep exploring the highest probability path
- Example
  - From node 0 to node 4. What is the best path?





# Example run

- At the beginning, we start at node 0 with probability of 1



Queue


| Node | Probability |
|------|-------------|
| 0    | 1           |

Visited

|  |
|--|
|  |
|--|

# Example run

- With no other choice in the queue, we check where we can go from node 0



Queue


| Node | Probability |
|------|-------------|
| 1    | 0.5         |
| 2    | 0.2         |
| 3    | 0.8         |

Visited

|   |
|---|
| 0 |
|---|

# Example run

- Next, we pick node 3 as it has the highest probability
  - Node 0 should **NOT** be added to the queue
    - Avoid infinite loop. Handled by using the visited **set**



Queue

| Node | Probability |
|------|-------------|
| 1    | 0.5         |
| 2    | 0.2         |
| 4    | 0.08        |

Visited


|      |
|------|
| 0, 3 |
|------|

Target is in the queue. Do we want to stop here?

Maybe there is another path, begin with lower probability but ends with a higher probability

# Example run

- Next, we pick node 1



Queue

| Node | Probability |
|------|-------------|
| 2    | 0.2         |
| 4    | 0.08        |
| 2    | 0.25        |
| 4    | 0.1         |

Visited


0, 3, 1

Two copies of node 4 in the queue. Reached by different paths. We only want the highest probability one.

Ignore the second copy and so on by using the visited

# Example run

- Next, we pick node 2 with probability of 0.25



Queue


| Node | Probability |
|------|-------------|
| 2    | 0.2         |
| 4    | 0.08        |
| 4    | 0.1         |
| 4    | 0.25        |

Visited

|         |
|---------|
| 0, 3, 1 |
|---------|

# Example run

- Next, we pick node 4 with probability of 0.25
  - This is the end. We reach the destination



Queue

| Node | Probability |
|------|-------------|
| 2    | 0.2         |
| 4    | 0.08        |
| 4    | 0.1         |

Visited

|         |
|---------|
| 0, 3, 1 |
|---------|

Q: There won't be another path reaching node 4 with a probability  $> 0.25$ . Why?

Q: What data structure should we use for the queue?

# Converting into code-like steps

- Pseudocode
  - Put (start, 1) to the queue
  - While queue is not empty
    - Get the record (u, p) with the highest probability
      - Recall: u is the node, p is the probability
    - If u is the end
      - Bingo. The answer is p
    - If u is not visited
      - Add u to visited
      - For all reachable nodes from u that are not visited
        - Add them to the queue
  - End of while loop
  - Still not finding the end? It means it is not reachable. Return 0

Check the sample codes to see how this is converted into Python

# Multiprocessing

- Run in one process vs run in multiple processes
  - If you write a normal Python program, it is only using one core of your CPU
- Theory and concepts behind multiprocessing will not be covered in this module

12th Gen Intel® Core™ Desktop Processors Comparison

|                                                                    | Intel® Core™ i9 | Intel® Core™ i7 | Intel® Core™ i5                          | Intel® Core™ i3 |
|--------------------------------------------------------------------|-----------------|-----------------|------------------------------------------|-----------------|
| Max Turbo Frequency [GHz]                                          | Up to 5.2       | Up to 5.0       | Up to 4.9                                | Up to 4.4       |
| Intel® Turbo Boost Max Technology 3.0 Frequency [GHz] <sup>1</sup> | Up to 5.2       | Up to 5.0       | n/a                                      | n/a             |
| Performance-core Max Turbo Frequency [GHz] <sup>6</sup>            | Up to 5.1       | Up to 4.9       | Up to 4.9                                | Up to 4.4       |
| Efficient-core Max Turbo Frequency [GHz] <sup>6</sup>              | Up to 3.9       | Up to 3.8       | Up to 3.6                                | n/a             |
| Performance-core Base Frequency [GHz]                              | Up to 3.2       | Up to 3.6       | Up to 3.7                                | Up to 3.5       |
| Efficient-core Base Frequency [GHz]                                | Up to 2.4       | Up to 2.7       | Up to 2.8                                | n/a             |
| Processor Cores (P-cores + E-cores) <sup>7</sup>                   | 16 (8P + 8E)    | 12 (8P + 4E)    | 10 (6P + 4E) or 6 (6P + 0E) <sup>7</sup> | 4 (4P + 0E)     |



# GPU vs CPU

- GPU is also processing unit like CPU
- GPU has much more cores than CPU, but the computational ability of each core is weaker than a core of CPU
- GPU is good for tasks that can be run in parallel
  - Like matrix operations (which is heavily used in neural network)
    - Do you see why we usually use GPU in AI?

# TPU? LPU?

- TPU – Tensor Processing Units
  - Customized to support computations around neural network (tensors)
    - Tensors: basic unit of data inside GPU (e.g., numbers / vectors / matrices)
  - Cons: not for other GPU tasks, e.g., video / image rendering
- LPU – Language Processing Units
  - Customized to support LLM operations

# Sample codes for multiprocessing

A multiprocessing executor in Python

```
executor = ProcessPoolExecutor()
pool = [executor.submit(a_long_task, id, id) for id in range(10)]

for task in as_completed(pool):
 print(task.result())
```

Submit the tasks to the executor

Get back the results

Note: the results may not be in the same order as they are submitted

# Data structure in pandas / numpy

- They are designed for supporting machine learning and data science operations, e.g., many operations are done on tabular data
- Check experiment 3

# Summary

- Remember, no single data structure is the best in all scenarios
- Analyze your use case and pick the best one
  - The performance difference can be **huge**, especially when the data is huge
- More advanced topics for self-study
  - RAM
  - SSD
  - File system